



# Carvana Image Masking Challenge: Implementing a U-net

Jiawei Yu

Feb. 22, 2019

Kaggle task link:

<https://www.kaggle.com/c/carvana-image-masking-challenge>

# Outline

- ◎ Task overview
- ◎ U-net
- ◎ Data preprocessing & generation
- ◎ Implementing & tweaking a U-net
- ◎ Data augmentation
- ◎ Results analyses & visualization



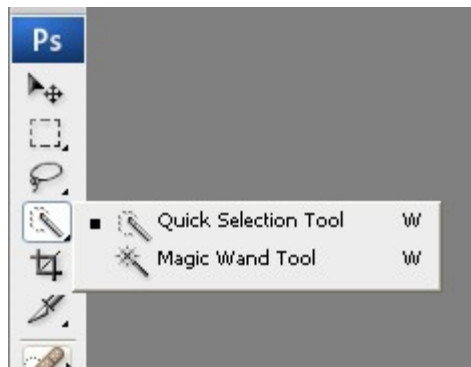


1.

# Task Overview

Training/testing data visualization,  
train/val/test split plan.

“

Develop an algorithm that **automatically** removes the photo studio background



# Training data

Image



## Image information


JPG format

RGB mode (3 channels)

'uint8' data ranging from 0 to 255

Size of  $1918 \times 1280$  pixels

Mask



## Mask information

GIF format

INDEX mode (1 channel)

'uint8' data with 0 and 1.

Size of  $1918 \times 1280$  pixels

## Train data w label (mask)

Total **5088** images/masks.

**318** unique cars, each has **16** images taken from 16 different angles.

## Test data w/o label

Total **100064** images/masks.

**6254** unique cars, each has **16** images taken from 16 different angles.




## Train data w label (mask)

Total **5088** images/masks.

**318** unique cars, each has **16** images taken from 16 different angles.

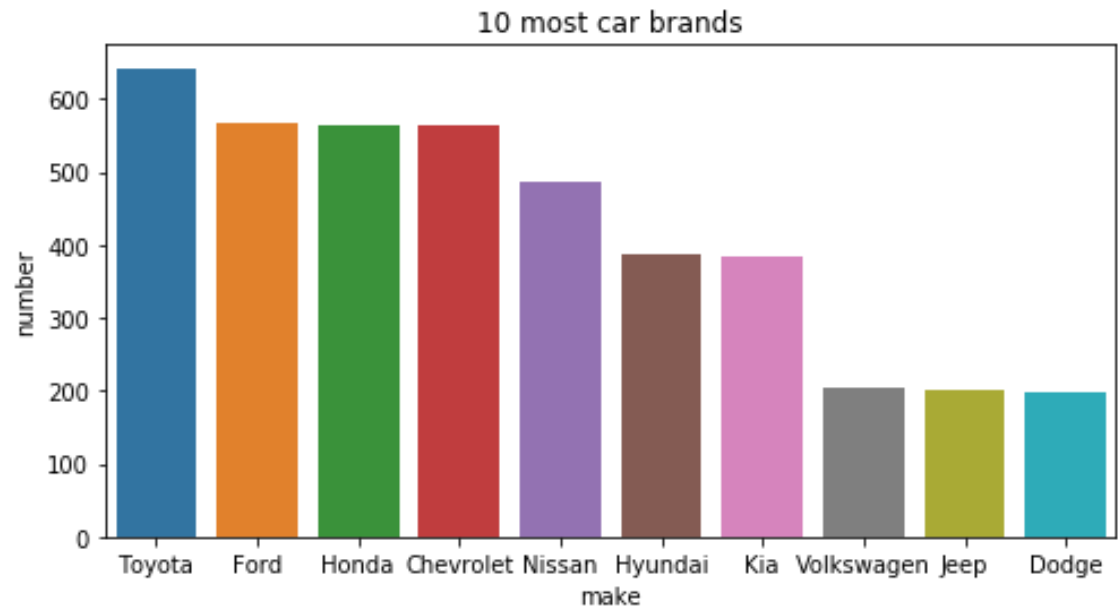
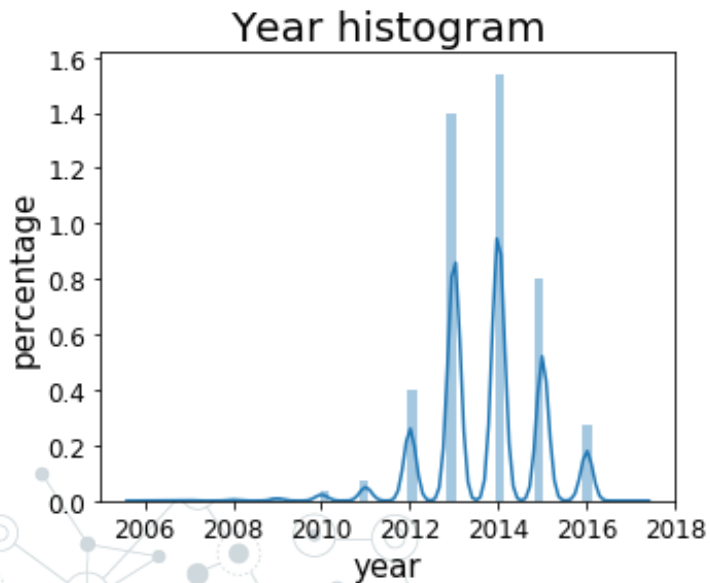


**Data leakage alert!** Images of a same car with different angles should NOT be split into different pools (train/validation/test)



# Meta data

	id	year	make	model	trim1	trim2
0	0004d4463b50	2014.0	Acura	TL	TL	w/SE
1	00087a6bd4dc	2014.0	Acura	RLX	RLX	w/Tech
2	000aa097d423	2012.0	Mazda	MAZDA6	MAZDA6	i Sport
3	000f19f6e7d4	2016.0	Chevrolet	Camaro	Camaro	SS
4	00144e887ae9	2015.0	Acura	TLX	TLX	SH-AWD V6 w/Advance Pkg








2.

# U-net

Powerful end-to-end CNN used for  
image segmentation

# U-net architecture






3.

# Data preprocessing & generation

Image resize, train/val/test split  
and data generator

# Resize the image & train/val/test split



```
dic_Im_names = {'train': Im_trains,  
                'val': Im_vals,  
                'test': Im_tests}
```




Contains all unique car names as a list.


# Data generator

```
1 print('X_train_raw shape is ', np.shape(X_train_raw))
2 print('y_train_raw shape is ', np.shape(y_train_raw))
3 print('X_train_raw size is ', X_train_raw.nbytes/1024/1024, 'Mb')
```

X\_train\_raw shape is (3040, 128, 128, 3)  
y\_train\_raw shape is (3040, 128, 128)  
X\_train\_raw size is 142.5 Mb



Large data size: use a data generator to grab and train the batch on the fly.





4.

# Implementing & tweaking a U-net

Tackling input size & neural network hyperparameters

# Loss function & metrics


## Loss function:

Binary cross entropy =  $-(y \log(p) + (1 - y) \log(1 - p))$

## Evaluation metric:

dice coefficient

$$DSC = \frac{2|X \cap Y|}{|X| + |Y|} = \frac{2 * y * round(p)}{y + round(p)}$$
$$= \frac{2TP}{2TP + FP + FN}$$

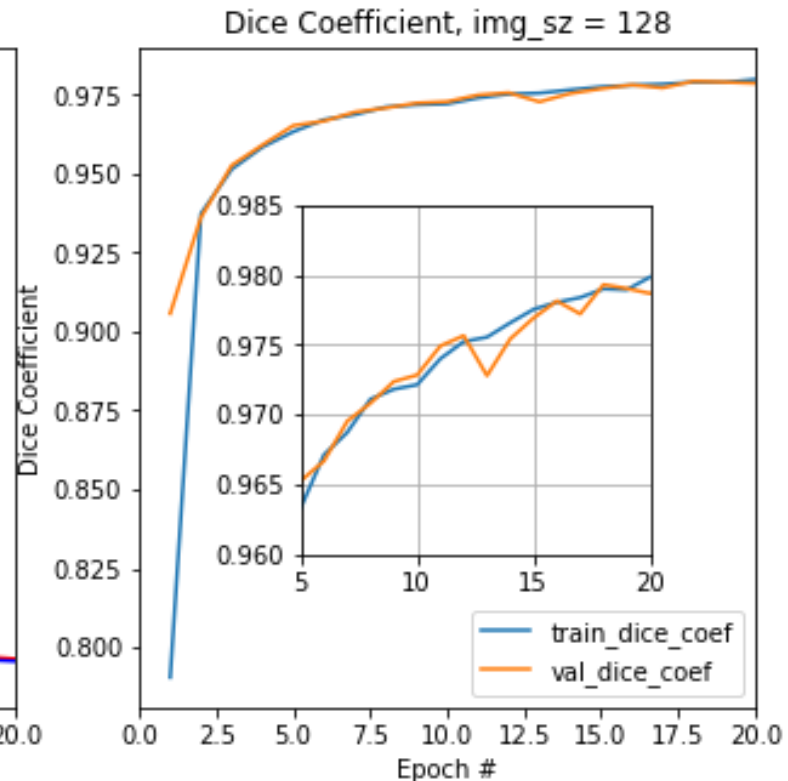
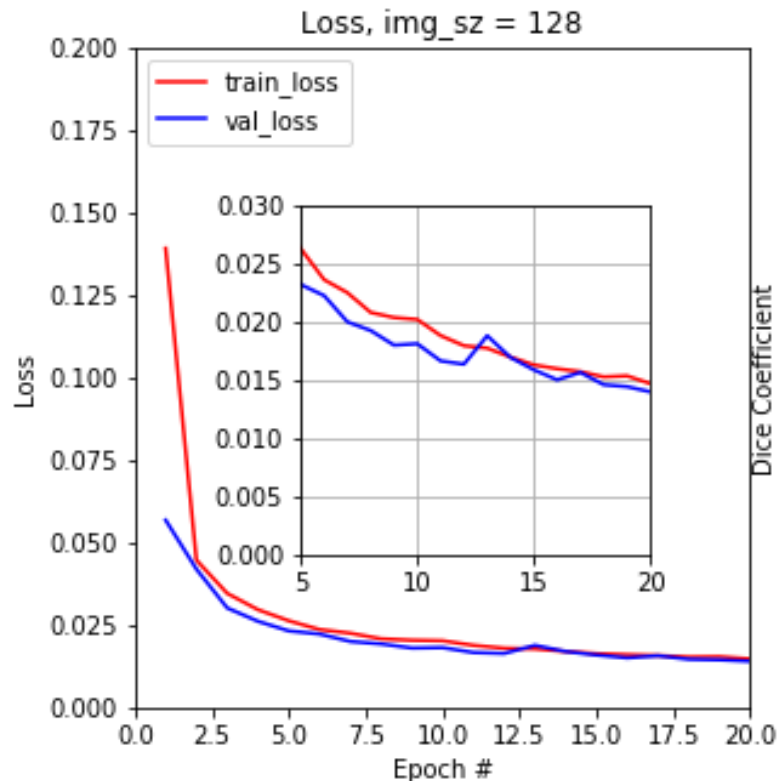


$$IoU = \frac{TP}{TP + FP + FN}$$

# Compile and train the model

Total params: 120,965  
Trainable params: 120,965  
Non-trainable params: 0


```
1 model.compile(optimizer=Adam(1e-3), loss='binary_crossentropy', metrics=[dice_coef])
2 history = model.fit(x = X_train, y = y_train,
3                     validation_data=(X_val, y_val),
4                     epochs = 30, batch_size = 16)
```



1024/1024 [=====] - 54s 53ms/step  
Test loss = 0.0167651942756  
Test dice\_coef = 0.976771984249



# Make a prediction


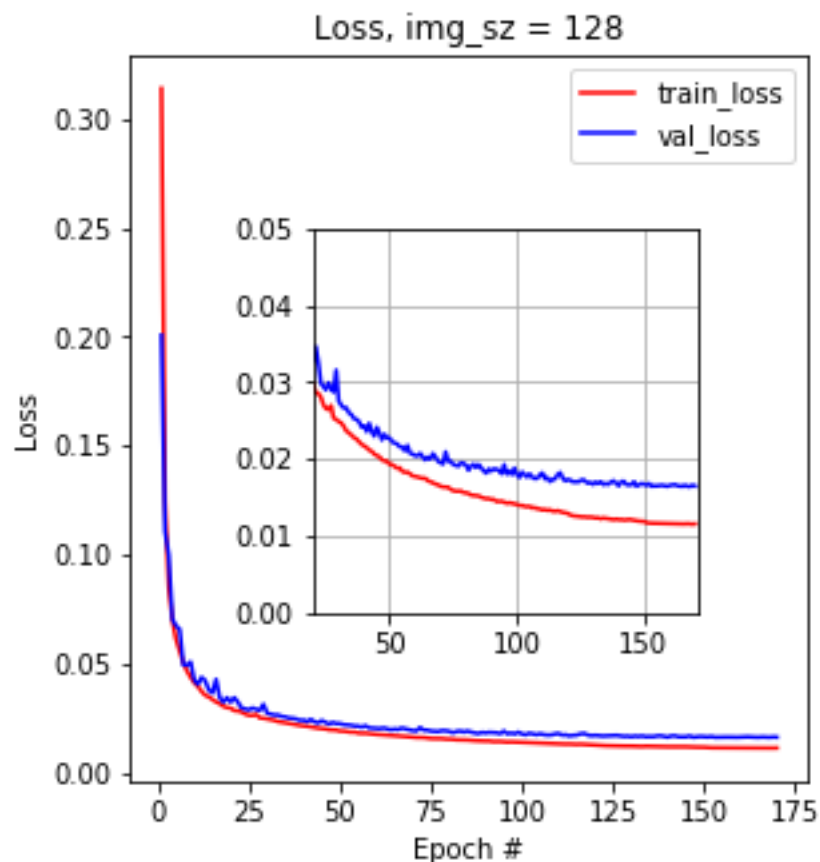


# Use a larger U-net

Total params: 485,957

Trainable params: 485,957

Non-trainable params: 0



# Increase the image size

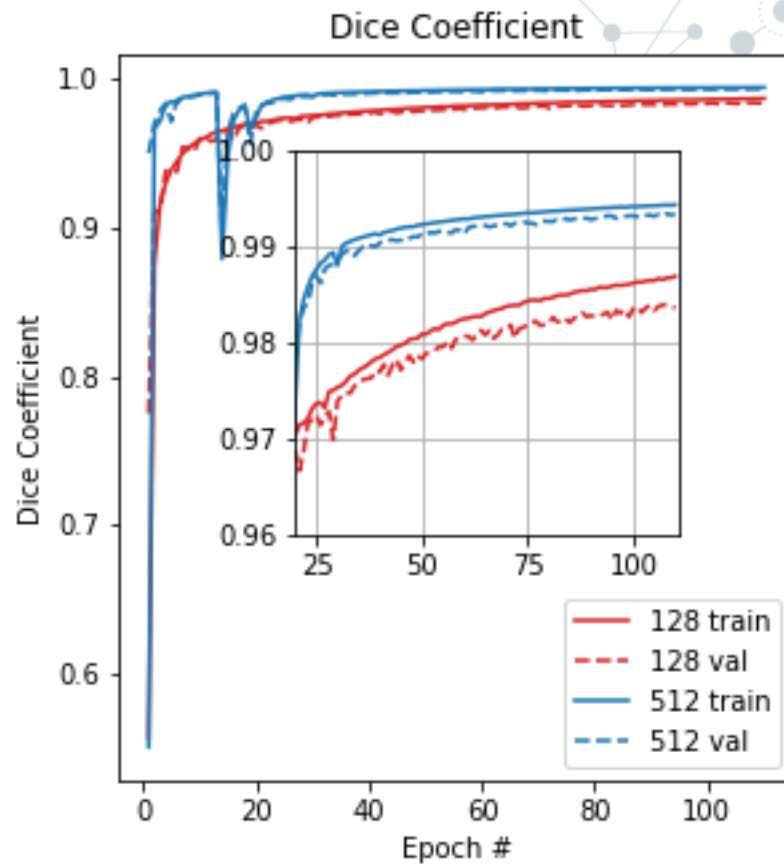
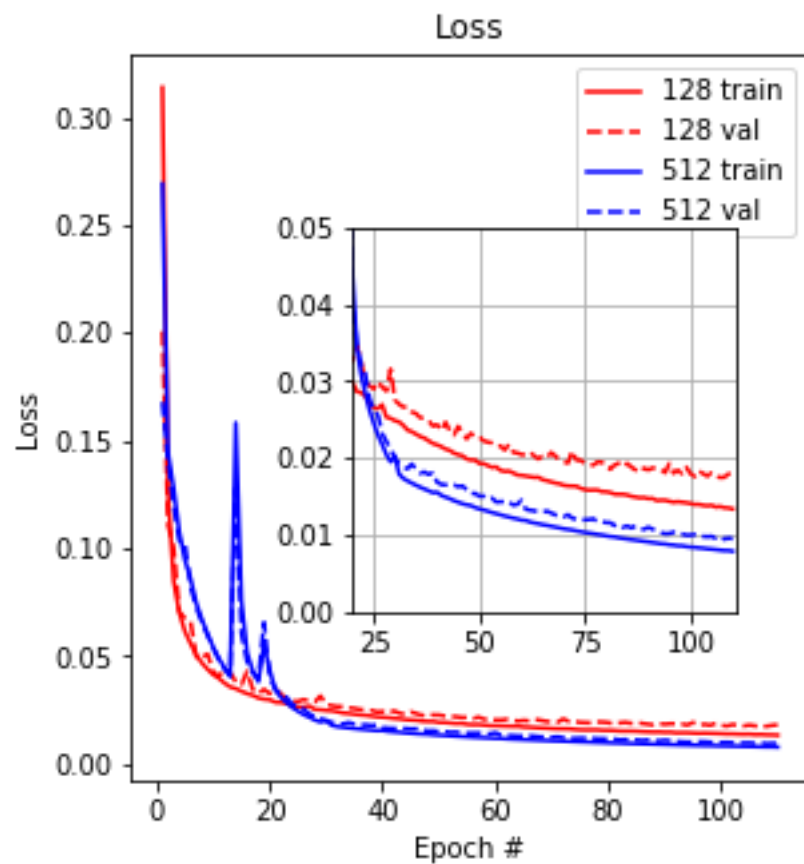


Image size = 512

```
1 score = model.evaluate_generator(test_generator)
2 print ("Test loss = ", score[0])
3 print ("Test dice_coef = ", score[1])
```

Test loss = 0.00861717724183

Test dice\_coef = 0.993911107071





5.

# Data augmentation

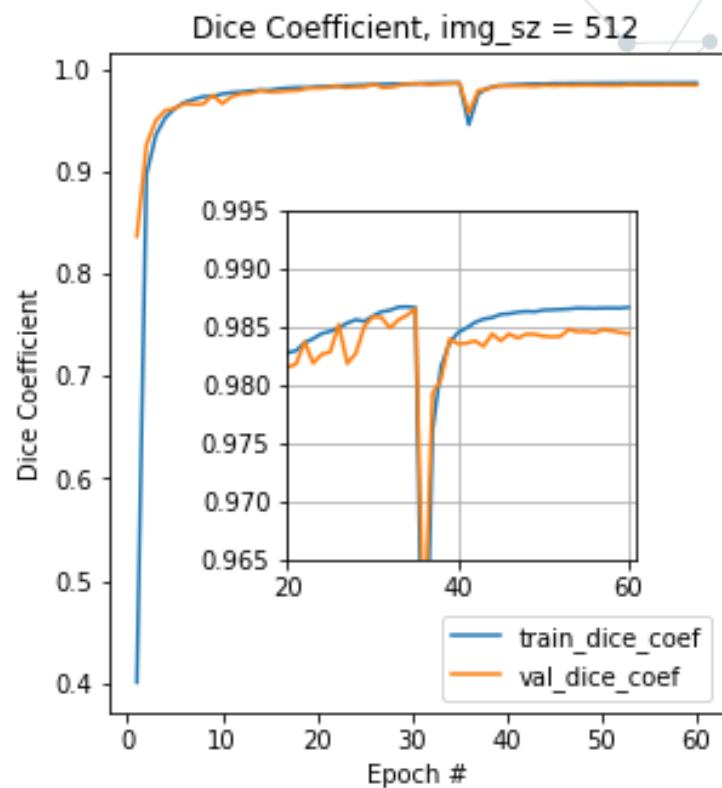

Generate more training data to  
train a larger U-net

# Random rotation & horizontal flip

- Random rotation ( $<25$  degree)
- 50% chance to be flipped horizontally.



Set mode='symmetric' to eliminate the black triangles.



Performance drops after implementing augmentation, indicating that **model with data augmentation underfits the data.**

train


aug\_off



aug\_on



test



Training images w/o augmentation share the same studio background with that of the test images. It will be easier for nn to learn these simple features and make a better prediction.

**Possible solution:** use **batch norm** to reduce the “covariance” of different layers in nn.




6.

# Results analyses & visualization

Model evaluation, results analyses and feature map visualization



# Model summary




Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	(None, 512, 512, 3)	0	
conv2d_49 (Conv2D)	(None, 512, 512, 8)	224	input_3[0][0]
conv2d_50 (Conv2D)	(None, 512, 512, 8)	584	conv2d_49[0][0]
max_pooling2d_9 (MaxPooling2D)	(None, 256, 256, 8)	0	conv2d_50[0][0]
conv2d_51 (Conv2D)	(None, 256, 256, 16)	1168	max_pooling2d_9[0][0]
conv2d_52 (Conv2D)	(None, 256, 256, 16)	2320	conv2d_51[0][0]
max_pooling2d_10 (MaxPooling2D)	(None, 128, 128, 16)	0	conv2d_52[0][0]
conv2d_53 (Conv2D)	(None, 128, 128, 32)	4640	max_pooling2d_10[0][0]
conv2d_54 (Conv2D)	(None, 128, 128, 32)	9248	conv2d_53[0][0]
max_pooling2d_11 (MaxPooling2D)	(None, 64, 64, 32)	0	conv2d_54[0][0]
conv2d_55 (Conv2D)	(None, 64, 64, 64)	18496	max_pooling2d_11[0][0]
conv2d_56 (Conv2D)	(None, 64, 64, 64)	36928	conv2d_55[0][0]
dropout_5 (Dropout)	(None, 64, 64, 64)	0	conv2d_56[0][0]
max_pooling2d_12 (MaxPooling2D)	(None, 32, 32, 64)	0	dropout_5[0][0]
conv2d_57 (Conv2D)	(None, 32, 32, 128)	73856	max_pooling2d_12[0][0]
conv2d_58 (Conv2D)	(None, 32, 32, 128)	147584	conv2d_57[0][0]
dropout_6 (Dropout)	(None, 32, 32, 128)	0	conv2d_58[0][0]
up_sampling2d_9 (UpSampling2D)	(None, 64, 64, 128)	0	dropout_6[0][0]
conv2d_59 (Conv2D)	(None, 64, 64, 64)	32832	up_sampling2d_9[0][0]
concatenate_9 (Concatenate)	(None, 64, 64, 128)	0	dropout_5[0][0] conv2d_59[0][0]
conv2d_60 (Conv2D)	(None, 64, 64, 64)	73792	concatenate_9[0][0]
conv2d_61 (Conv2D)	(None, 64, 64, 64)	36928	conv2d_60[0][0]
up_sampling2d_10 (UpSampling2D)	(None, 128, 128, 64)	0	conv2d_61[0][0]
conv2d_62 (Conv2D)	(None, 128, 128, 32)	8224	up_sampling2d_10[0][0]
concatenate_10 (Concatenate)	(None, 128, 128, 64)	0	conv2d_54[0][0] conv2d_62[0][0]
conv2d_63 (Conv2D)	(None, 128, 128, 32)	18464	concatenate_10[0][0]
conv2d_64 (Conv2D)	(None, 128, 128, 32)	9248	conv2d_63[0][0]
up_sampling2d_11 (UpSampling2D)	(None, 256, 256, 32)	0	conv2d_64[0][0]
conv2d_65 (Conv2D)	(None, 256, 256, 16)	2064	up_sampling2d_11[0][0]
concatenate_11 (Concatenate)	(None, 256, 256, 32)	0	conv2d_52[0][0] conv2d_65[0][0]
conv2d_66 (Conv2D)	(None, 256, 256, 16)	4624	concatenate_11[0][0]
conv2d_67 (Conv2D)	(None, 256, 256, 16)	2320	conv2d_66[0][0]
up_sampling2d_12 (UpSampling2D)	(None, 512, 512, 16)	0	conv2d_67[0][0]
conv2d_68 (Conv2D)	(None, 512, 512, 8)	520	up_sampling2d_12[0][0]
concatenate_12 (Concatenate)	(None, 512, 512, 16)	0	conv2d_50[0][0] conv2d_68[0][0]
conv2d_69 (Conv2D)	(None, 512, 512, 8)	1160	concatenate_12[0][0]
conv2d_70 (Conv2D)	(None, 512, 512, 8)	584	conv2d_69[0][0]
conv2d_71 (Conv2D)	(None, 512, 512, 2)	146	conv2d_70[0][0]
conv2d_72 (Conv2D)	(None, 512, 512, 1)	3	conv2d_71[0][0]

Total params: 485,957  
Trainable params: 485,957  
Non-trainable params: 0


# Model evaluation: Good predictions

	Im_names_test	scores
113	2267f4aa0d2c_02.jpg	0.997615
114	2267f4aa0d2c_03.jpg	0.997280
502	9ab2a45de8c7_07.jpg	0.997027
127	2267f4aa0d2c_16.jpg	0.996976
126	2267f4aa0d2c_15.jpg	0.996976

Input image



Target mask



Predicted mask




# Bad predictions

	Im_names_test	scores
<b>104</b>	1390696b70b6_09.jpg	0.982636
<b>256</b>	ae296a20fdd9_01.jpg	0.983955
<b>103</b>	1390696b70b6_08.jpg	0.984998
<b>471</b>	61060ada97c9_08.jpg	0.985197
<b>920</b>	28109f18d9d4_09.jpg	0.985368

Input image



Target mask



Predicted mask



	lm_names_test	scores
256	ae296a20fdd9_01.jpg	0.983955
103	1390696b70b6_08.jpg	0.984998

Input image




Target mask



Predicted mask




Input image



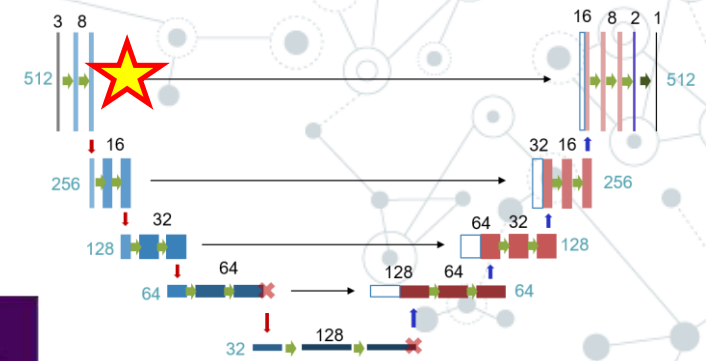
Target mask

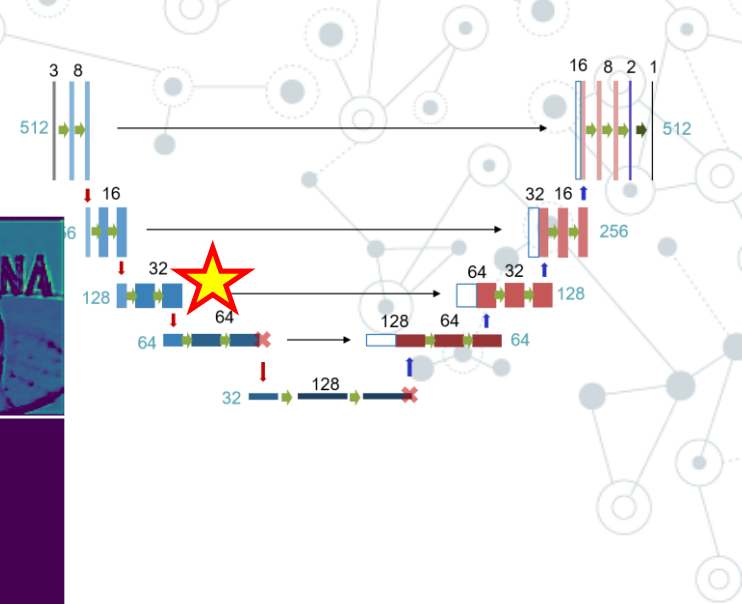



Predicted mask



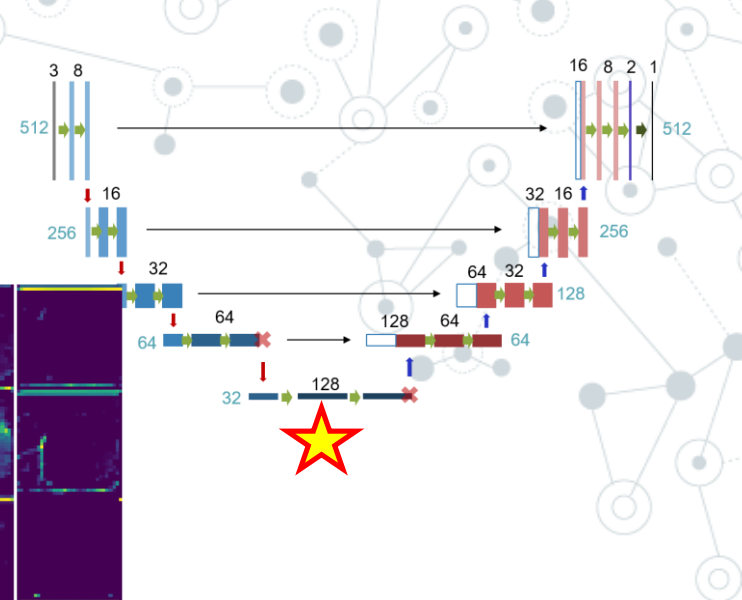

# Visualization of Feature map



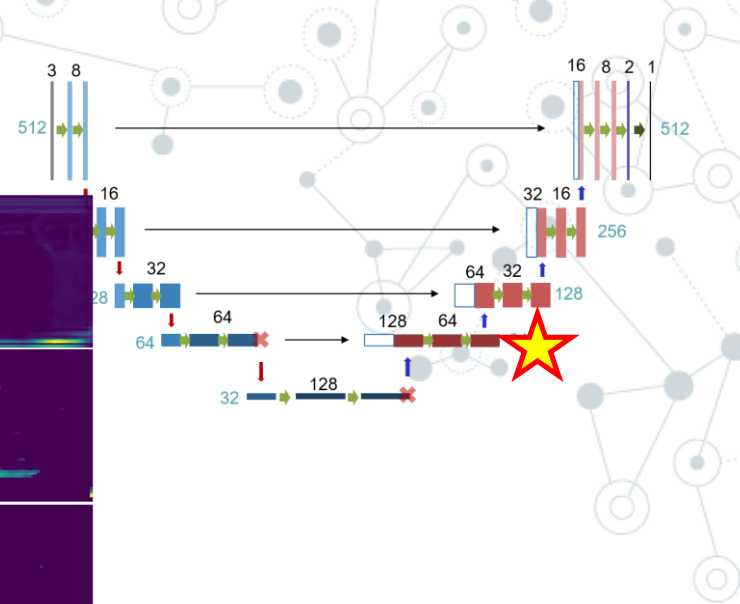



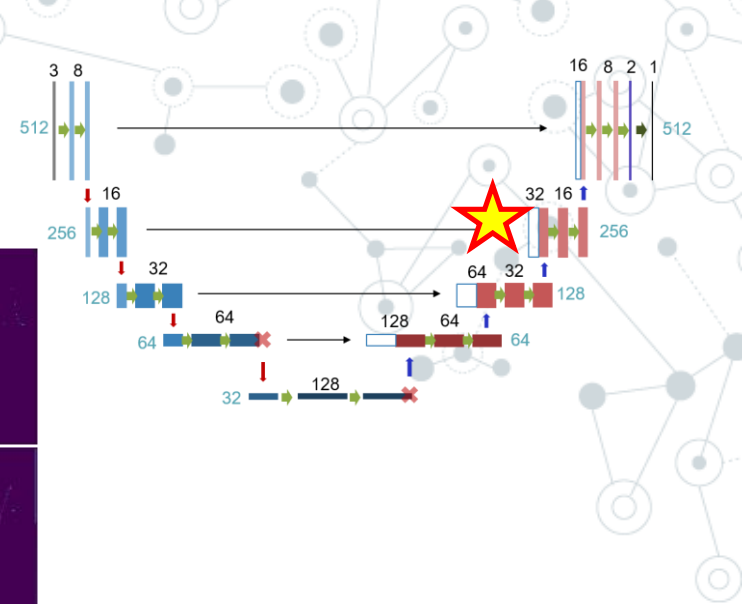



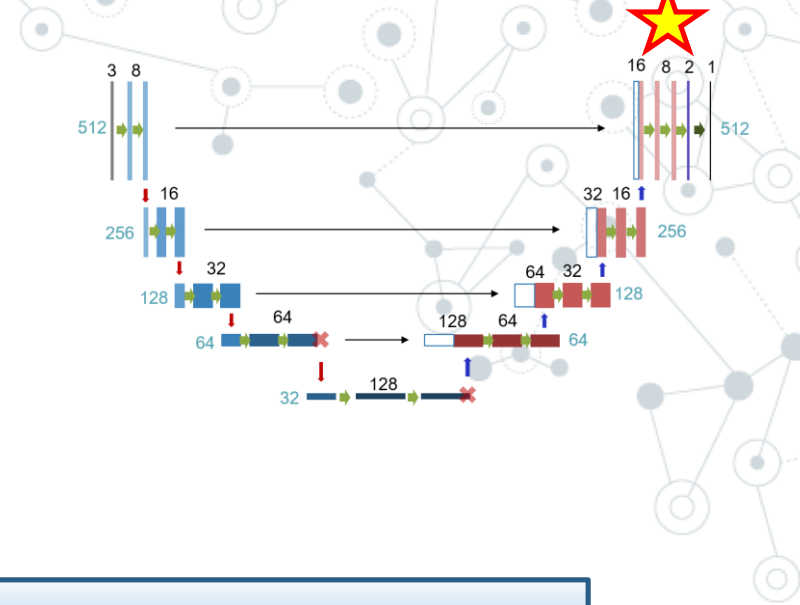












And finally...a prediction!



# Here is what I learnt...



## U-net

A powerful CNN architecture for semantic segmentation tasks.



## Data Generator

Train NN on the fly. Practiced OOP.



## Keras

Awesome API to implement CNN in a few lines of codes.



## Data Augmentation

Generate enormous data to train a large NN.



## Visualization

Plot results for a better analyse and evaluation.



## Feature map

"Look into" the deep NN and see what it is learning.

# Thanks!

**Any questions?**






# Appendix

Some useful stuff..




# Upsampling: interpolation




# Transposed convolution

Conv



Trans Conv



Convolution By Matrix Multiplication

Source: <https://towardsdatascience.com/up-sampling-with-transposed-convolution-9ae4f2df52d0>



# Transposed convolution

